



Larger models do not solve the small spaceflight training-set problem

Current single-cell foundation models and text LLMs do not reliably beat a tuned classical baseline on these held-out mission tasks.

A. Classical baseline is strongest across the shared 6-task set



B. Foundation-model gains are local; most tissues are lower



C. Evaluation scope

